

Distributed Scheduling

Distributed Scheduling


Goal: enable transparent execution of programs on networked computing systems

Motivations: reduce response time of program execution through load balancing


Distributed Scheduling

- Let P be the probability that the system is in a state in which at least 1 task is waiting for service and at least 1 server is idle.
- Let ρ be the utilization of each server.
- We can estimate P using probabilistic analysis and plot a graph against system utilization.

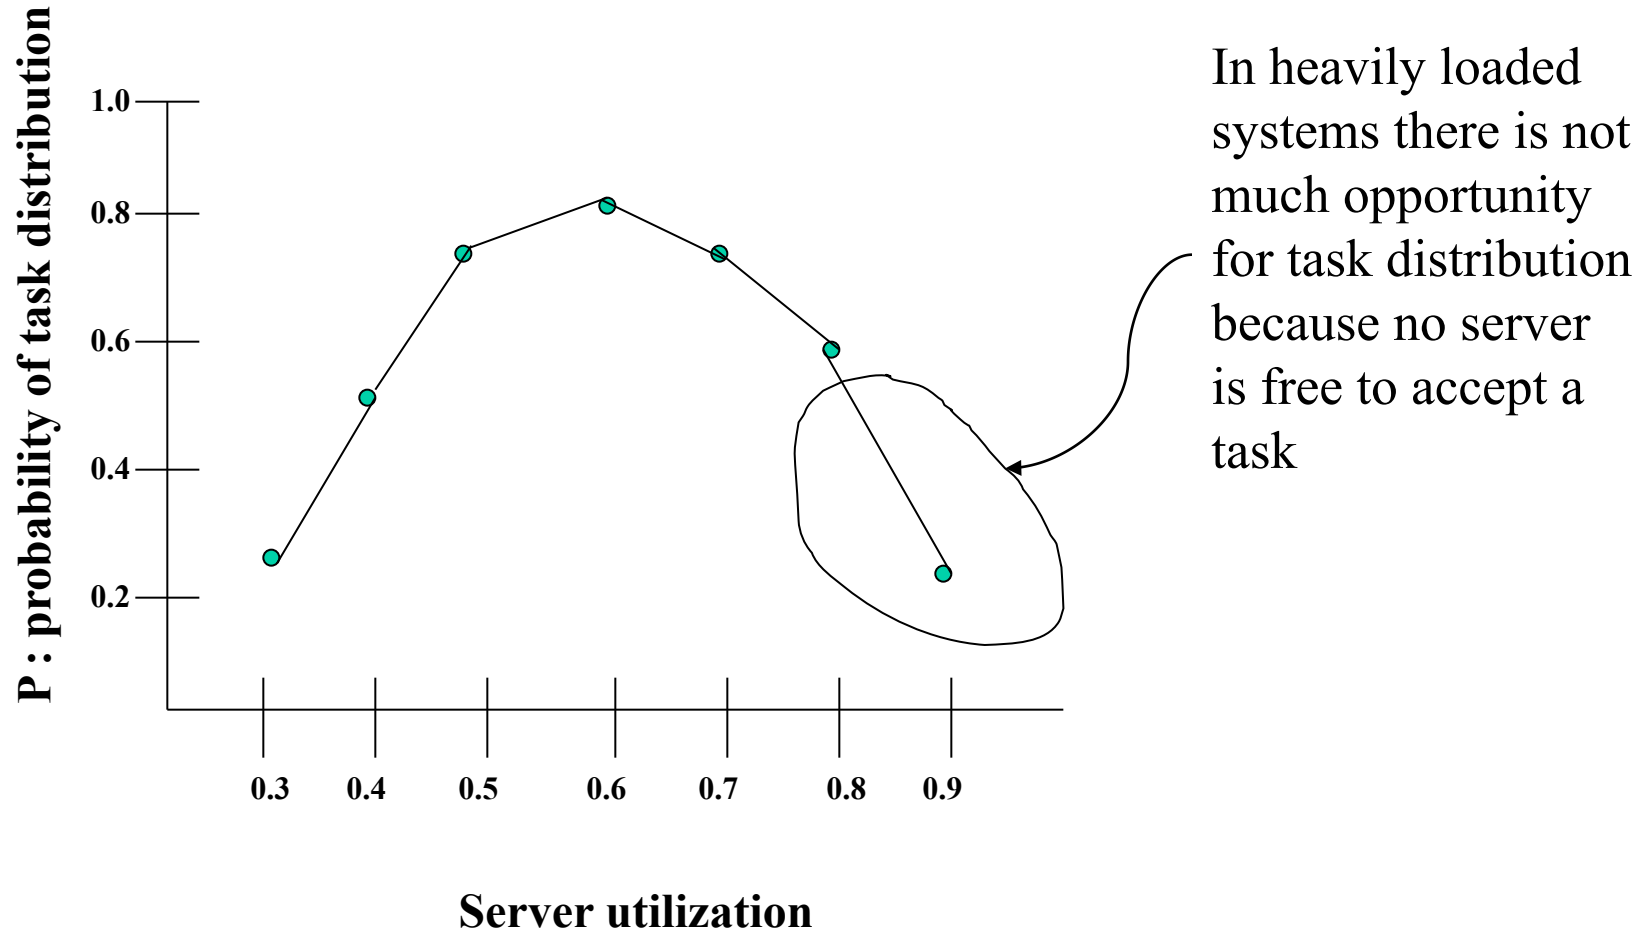
Opportunities for Task Distribution



Task Distribution




Task Distribution



Task Distribution

P : probability of task distribution



Server utilization

In moderately loaded systems there are good opportunities to distribute tasks from over-utilized to under-utilized systems

What is Load?

- Load on a system/node can correspond to the queue length of tasks/ processes that need to be processed.
- Queue length of waiting tasks: proportional to task response time, hence a good indicator of system load.
- Distributing load: transfer tasks/processes among nodes.
- If a task transfer (from another node) takes a long time, the node may accept more tasks during the transfer time.
- Causes the node to be highly loaded. Affects performance.

Types of Algorithms

- *Static load distribution algorithms*: Decisions are hard-coded into an algorithm with a priori knowledge of system.
- *Dynamic load distribution*: use system state information such as task queue length, processor utilization.
- *Adaptive load distribution*: adapt the approach based on system state.
 - (e.g.,) Dynamic distribution algorithms collect load information from nodes even at very high system loads.
 - Load information collection itself can add load on the system as messages need to be exchanged.
 - Adaptive distribution algorithms may stop collecting state information at high loads.

Balancing vs. Sharing

- *Load balancing*: Equalize load on the participating nodes.
 - Transfer tasks even if a node is not heavily loaded so that queue lengths on all nodes are approximately equal.
 - More number of task transfers, might degrade performance.
- *Load sharing*: Reduce burden of an overloaded node.
 - Transfer tasks only when the queue length exceeds a certain threshold.
 - Less number of task transfers.
- *Anticipatory task transfers*: transfer from overloaded nodes to ones that are likely to become idle/lightly loaded.
 - More like load balancing, but may be less number of transfers.

Types of Task Transfers

- *Preemptive task transfers*: transfer tasks that are partially executed.
 - Expensive as it involves collection of task states.
 - Task state: virtual memory image, process control block, IO buffers, file pointers, timers, ...
- *Non-preemptive task transfers*: transfer tasks that have not begun execution.
 - Do not require transfer of task states.
 - Can be considered as task placements. Suitable for load sharing not for load balancing.
- Both transfers involve information on user's current working directory, task privileges/priority.

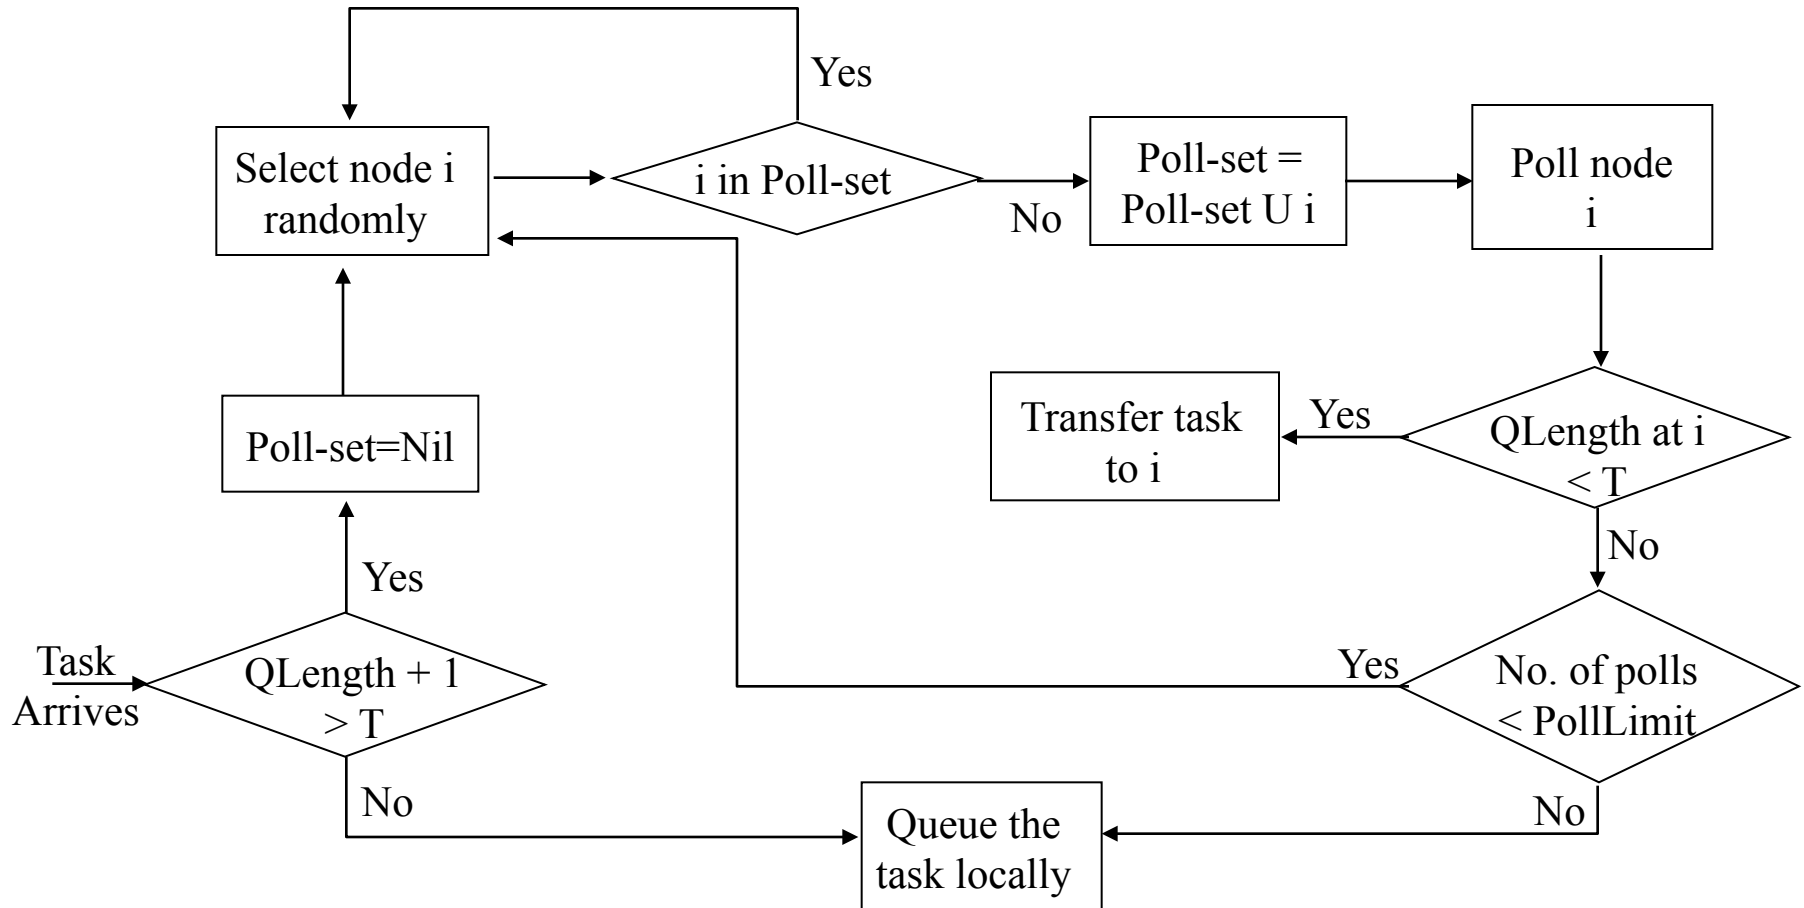
Load Distributing Algorithms

- Sender-initiated: distribution initiated by an overloaded node.
- Receiver-initiated: distribution initiated by lightly loaded nodes.
- Symmetric: initiated by both senders and receivers. Has advantages and disadvantages of both the approaches.
- Adaptive: sensitive to state of the system.

Sender-initiated

- *Transfer Policy*: Use thresholds.
 - Sender if queue length exceeds T .
 - Receiver if accepting a task will not make queue length exceed T .
- *Selection Policy*: Only newly arrived tasks.
- *Location Policy*:
 - Random: Use no remote state information. Task transferred to a node at random.
 - No need for state collection. Unnecessary task transfers (processor thrashing) may occur.
 - Threshold: poll a node to find out if it is a receiver. Receiver must accept the task irrespective of when it (task) actually arrives.
 - *PollLimit*, ie., the number of polls, can be used to reduce overhead.
 - Shortest: Poll a set of nodes. Select the receiver with shortest task queue length.

Sender-initiated



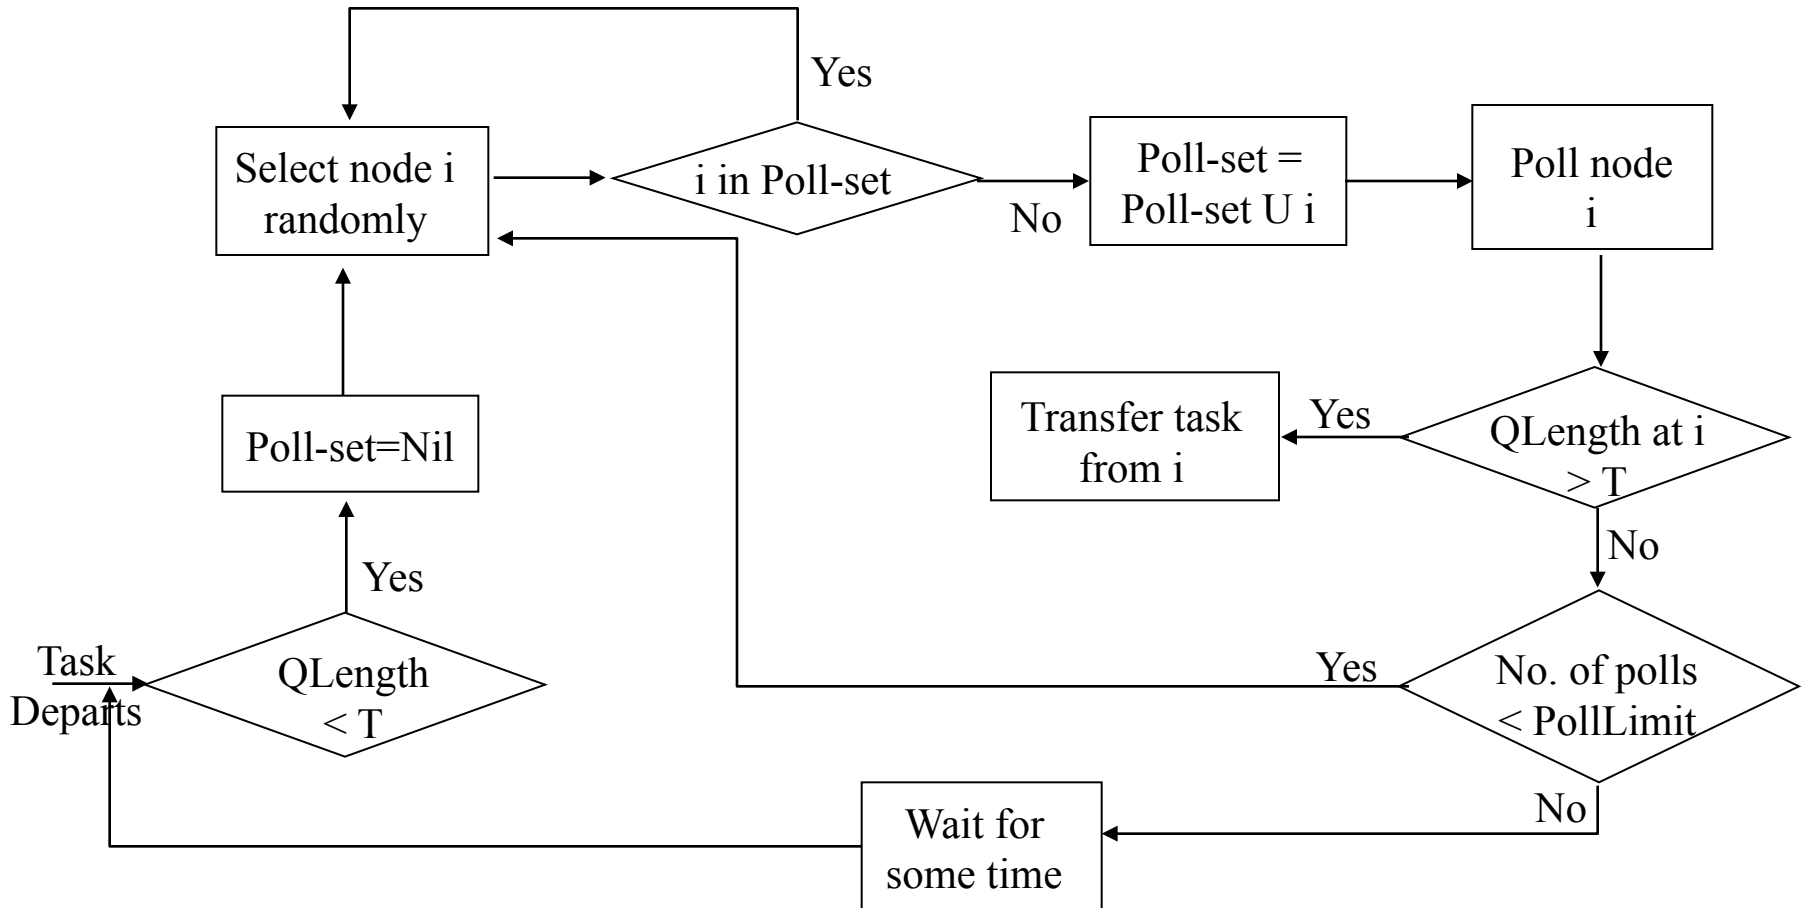
Sender-initiated

- *Information Policy*: demand-driven.
- *Stability*: can become unstable at high loads.
 - At high loads, it may become difficult for senders to find receivers.
 - Also, the number of senders increase at high system loads thereby increasing the polling activity.
 - Polling activity may make the system unstable at high loads.

Receiver-initiated

- *Transfer Policy*: uses thresholds. Queue lengths below T identifies receivers and those above T identifies senders.
- *Selection Policy*: as before.
- *Location Policy*: Polling.
 - A random node is polled to check if a task transfer would place its queue length below a threshold.
 - If not, the polled node transfers a task.
 - Otherwise, poll another node till a static PollLimit is reached.
 - If all polls fail, wait until another task is completed before starting polling operation.
- Information policy: demand-driven.
- Stability: Not unstable since there are lightly loaded systems that have initiated the algorithm.

Receiver-initiated



Receiver-initiated

- Drawback:
 - Polling initiated by receiver implies that it is difficult to find senders with new tasks.
 - Reason: systems try to schedule tasks as and when they arrive.
 - Effect: receiver-initiated approach might result in preemptive transfers. Hence transfer costs are more.
 - Sender-initiated: transfer costs are low as new jobs are transferred and so no need for transferring task states.

Symmetric

- Senders search for receivers and vice-versa.
- Low loads: senders can find receivers easily. High loads: receivers can find senders easily.
- May have disadvantages of both: polling at high loads can make the system unstable. Receiver-initiated task transfers can be preemptive and so expensive.
- Simple algorithm: combine previous two approaches.
- Above-average algorithm:
 - Transfer Policy: Two adaptive thresholds instead of one. If a node's estimated average load is A , a higher threshold *TooHigh* $> A$ and a lower threshold *TooLow* $< A$ are used.
 - Load $< \text{TooLow}$ \rightarrow receiver. Load $> \text{TooHigh}$ \rightarrow sender.

Above-average Algorithm

- Location policy:
- Sender Component
 - Node with *TooHigh* load, broadcasts a *TooHigh* message, sets *TooHigh* timer, and listens for an *Accept* message.
 - A receiver that gets the (TooHigh) message sends an *Accept* message, increases its load, and sets *AwaitingTask* timer.
 - If the *AwaitingTask* timer expires, load is decremented.
 - On receiving the *Accept* message: if the node is still a sender, it chooses the best task to transfer and transfers it to the node.
 - When sender is waiting for *Accept*, it may receive a *TooLow* message (receiver initiated). Sender sends *TooHigh* to that receiver. Do step 2 & 3.
 - On expiration of *TooHigh* timer, if no *Accept* message is received, system is highly loaded. Sender broadcasts a *ChangeAverage* message.

Above-average Algorithm...

- Receiver Component
 - Node with *TooLow* load, broadcasts a *TooLow* message, sets a *TooLow* timer, and listens for *TooHigh* message.
 - If *TooHigh* message is received, do step 2 & 3 in Sender Component.
 - If *TooLow* timer expires before receiving any *TooHigh* message, receiver broadcasts a *ChangeAverage* message to decrease the load estimate at other nodes.
- Selection Policy: as discussed before.
- Information policy: demand driven. Average load is modified based on system load. High loads may have less number of senders progressively.
 - Average system load is determined individually. There is a range of acceptable load before trying to be a sender or a receiver.

Adaptive Algorithms

- Limit Sender's polling actions at high load to avoid instability.
- Utilize the collected state information during previous polling operations to classify nodes as: Sender/overloaded, receiver/underloaded, OK (in acceptable load range).
- Maintained as separate lists for each class.
- Initially, each node assumes that all others are receivers.
- Location policy at sender:
 - Sender polls the head of the receiver list.
 - Polled node puts the sender at the head of its sender list. It informs the sender whether it is a receiver, a sender, or a OK node.
 - If the polled node is still a receiver, the new task is transferred.
 - Else the sender updates the polled node's status, polls the next potential receiver.
 - If this polling process fails to identify a receiver, the task can still be transferred during a receiver-initiated dialogue.

Adaptive Algorithms...

- Location policy at receiver
 - Receivers obtain tasks from potential senders. Lists are scanned in the following order.
 - Head to tail in senders list (most up-to-date info used), tail to head in OK list (least up-to-date used), tail to head in receiver list.
 - Least up-to-date used in the hope that status might have changed.
 - Receiver polls the selected node. If the node is a sender, a task is transferred.
 - If the node is not a sender, both the polled node and receiver update each other's status.
 - Polling process stops if a sender is found or a static PollLimit is reached.


Adaptive Algorithms...

- At high loads, sender-initiated polling gradually reduces as nodes get removed from receiver list (and become senders).
 - Whereas at low loads, sender will generally find some receiver.
- At high loads, receiver-initiated works and can find a sender.
 - At low loads, receiver may not find senders, but that does not affect the performance.
- Algorithm dynamically becomes sender-initiated at low loads and receiver-initiated at high loads.
- Hence, algorithm is stable and can use non-preemptive transfers at low loads (sender initiated).

Selecting an Algorithm

- If a system never gets highly loaded, sender-initiated algorithms work better.
- Stable, receiver-initiated algorithms better for high loads.
- Widely fluctuating loads: stable, symmetric algorithms.
- Widely fluctuating loads + high migration cost for preemptive transfers: stable, sender-initiated algorithms.
- Heterogeneous work arrival: stable, adaptive algorithms.

Performance Comparison



Implementation Issues

- Task placement
 - A task that is yet to begin is transferred to a remote machine, and starts its execution there.
- Task migration
 - *State Transfer*:
 - State includes contents of registers, task stack, task status (running, blocked, etc.), file descriptors, virtual memory address space, temporary files, buffered messages.
 - Other info: current working directory, signal masks and handlers, references to children processes.
 - Task is frozen (suspended) before state transfer.
 - *Unfreeze*:
 - Task is installed at the new machine, unfrozen, and is put in the ready queue.

State Transfer

- Issues to be considered:
 - Cost to support remote execution including delays due to freezing the task.
 - Duration of freezing the task should be small. Otherwise it can result in timeouts of tasks interacting with the frozen one.
 - Transfer memory as they are referenced to reduce delay?
 - Residual Dependencies: amount of resources to be dedicated at the former host for the migrated task.
 - An implementation that does not transfer all the virtual memory address space at the time of migration, but transfers them as they are referenced.
 - Need for redirection of messages to the migrated task.
 - Location-dependent system calls at the former node.

State Transfer...

- Disadvantages of residual dependencies :
 - Affects reliability as it becomes dependent on former host
 - Affects performance since memory accesses can become slow as pages need to be transferred
 - Affects complexity as task states are distributed on several hosts.
- Location transparency:
 - Migration should hide the location of tasks
 - Message passing, process naming, file handling, and other activities should be transparent of actual location.
 - Task names and their locations can be maintained as hints. If the hints fail, they can be updated by a broadcast query or through a name server.