

# Time and synchronization

(“There’s never enough time...”)


# Why Global Timing?

- Suppose there were a globally consistent time standard
- Would be handy
  - Who got last seat on airplane?
  - Who submitted final auction bid before deadline?
  - Did defense move before snap?

# Time Standards

- UT1
  - Based on astronomical observations
  - “Greenwich Mean Time”
- TAI
  - Started Jan 1, 1958
  - Each second is 9,192,631,770 cycles of radiation emitted by Cesium atom
  - Has diverged from UT1 due to slowing of earth's rotation
- UTC
  - TAI + leap seconds to be within 800ms of UT1
  - Currently 35
  - Most recent: June 30, 2012

# Comparing Time Standards



# Clocks

- Piezoelectric effect:
  - Squeeze a quartz crystal:  
generates electric field
  - Apply electric field: crystal bends
- Quartz crystal clock:
  - Resonation like a tuning fork
  - Accurate to parts per million
  - Gain/lose  $\frac{1}{2}$  second per day

# Challenges

- Two clocks do not agree perfectly
- **Skew:** The time difference between two clocks
- Quartz oscillators vibrate at different rates
- **Drift:** The difference in rates of two clocks
- If we had two perfect clocks:
  - Skew = 0
  - Drift = 0

# When we detect a clock has a skew

- Eg: it is 5 seconds behind
- Or 5 seconds ahead
- What can we do?

# When we detect a clock has a skew

- e.g. it is 5 seconds behind
  - We can advance it 5 seconds to correct
  - Might skip over event scheduled in-between
- Or 5 seconds ahead
  - Pushing back 5 seconds is a bad idea
    - Message was received before it was sent
    - Document closed before it was saved etc..
  - We want **monotonicity**: time always increases
  - We want **continuity**: time doesn't make jumps



# When we detect a clock has a skew

- e.g. it is behind
  - Run it faster until it catches up
- It is ahead
  - Run it slower until it catches up
- This does not guarantee correct clock in future
  - Need to check and adjust periodically

# Distributed time

- Premise
  - The notion of time is well-defined (and measurable) at each single location
  - But the relationship between time at different locations is unclear
    - Can minimize discrepancies, but never eliminate them
- Reality
  - Stationary GPS receivers can get global time with  $< 1\mu\text{s}$  error
  - Few systems designed to use this

# A football example

- Five locations: Goal Keeper K1, Player P1, Player P2, Player P3 and Goal Keeper K2
- Ten events:
  - $e_1$ : keeper K1 kicks the ball
  - $e_2$ : ball arrives to Player P1
  - $e_3$ : P1 kicks the ball to Player P2
  - $e_4$ : P2 runs with the ball
  - $e_5$ : P2 kicks the ball towards Player P3
  - $e_6$ : P3 kicks the ball towards other keeper K2
  - $e_7$ : The ball arrives near the keeper
  - $e_8$ : But another Player P4 touches the ball
  - $e_9$ : Ball also touches the keeper K2
  - $e_{10}$ : Ball went inside the goal

# A football example

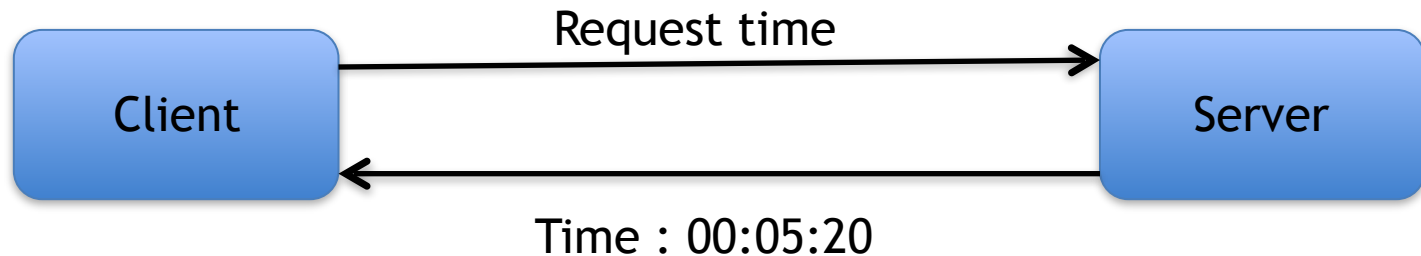
- Players and umpire knows  $e_1$  happens before  $e_6$ , which happens before  $e_7$
- Umpire knows  $e_2$  is before  $e_3$ , which is before  $e_4$ , which is before  $e_8$ , ...
- Relationship between  $e_8$  and  $e_9$  is unclear

# Ways to synchronize

- Send message from a player to another?
  - Or to a central timekeeper
  - How long does this message take to arrive?
- Synchronize clocks before the game?
  - Clocks drift
    - million to one => 1 second in 11 days
- Synchronize continuously during the game?
  - GPS, pulsars, etc


# How clocks synchronise

- Obtain time from time server:



# Perfect networks


- Messages always arrive, with propagation delay exactly  $d$



- Sender sends time  $T$  in a message
- Receiver sets clock to  $T+d$ 
  - Synchronization is exact

# How clocks synchronise

- Obtain time from time server:



- Time is inaccurate
  - Delays in message transmission
  - Delays due to processing time
  - Server's time may be inaccurate




# Synchronization in the real world

- Real networks are asynchronous
  - Propagation delays are arbitrary
- Real networks are unreliable
  - Messages don't always arrive

# Cristian's algorithm

- Request time, get reply
  - Measure actual round-trip time  $d$




- Sender's time was  $T$  between  $t_1$  and  $t_2$
- Receiver sets time to  $T + d/2$ 
  - Synchronization error is at most  $d/2$
- Can retry until we get a relatively small  $d$

# The Berkeley algorithm

- Master uses Cristian's algorithm to get time from many clients
  - Computes average time
  - Can discard outliers
- Sends time adjustments back to all clients

# The Network Time Protocol (NTP)

- Uses a hierarchy of time servers
  - Class 1 servers have highly-accurate clocks
    - connected directly to atomic clocks, etc.
  - Class 2 servers get time from only Class 1 and Class 2 servers
  - Class 3 servers get time from any server
- Synchronization similar to Cristian's alg.
  - Modified to use multiple one-way messages instead of immediate round-trip
- Accuracy: Local ~1ms, Global ~10ms



Source: [http://upload.wikimedia.org/wikipedia/commons/c/c9/Network\\_Time\\_Protocol\\_servers\\_and\\_clients.svg](http://upload.wikimedia.org/wikipedia/commons/c/c9/Network_Time_Protocol_servers_and_clients.svg)

# Real synchronization is imperfect

- Clocks never exactly synchronized
- Often inadequate for distributed systems
  - might need totally-ordered events
  - might need millionth-of-a-second precision

# Logical clocks

- Why do we need clocks?
  - To determine when one thing happened before another
- Can we determine that without using a “clock” at all?
  - Then we don’t need to worry about synchronisation, millisecond errors etc..


# Happened before

- $a \longrightarrow b$  : a happened before b
  - If a and b are successive events in same process then  $a \longrightarrow b$
  - Send before receive
    - If a : “send” event of message m
    - And b : “receive” event of message m
    - Then  $a \longrightarrow b$
  - Transitive:  $a \longrightarrow b$  and  $b \longrightarrow c \implies a \longrightarrow c$




# Example

- Events without a happened before relation are “concurrent”
- $e1 \longrightarrow e2$ ,  $e3 \longrightarrow e4$ ,  $e1 \longrightarrow e5$ ,  $e5 \parallel e2$



# Example

- Events without a happened before relation are “concurrent”
- Happened before is a partial ordering



# Happened before & causal order

- Happened before ==  
could have caused/influenced
- Preserves causal relations
- Implies a partial order
  - Implies time ordering between certain pairs of events
  - Does not imply anything about ordering between concurrent events


# Logical clocks

- Idea: Use a counter at each process
- Increment after each event
- Can also increment when there are no events
  - Eg. A clock
- An actual clock can be thought of as such an event counter
- It counts the states of the process
- Each event has an associated time: The count of the state when the event happened

# Lamport clocks

- Keep a logical clock (counter)
- Send it with every message
- On receiving a message, set own clock to  $\max(\{\text{own counter, message counter}\}) + 1$
- For any event  $e$ , write  $c(e)$  for the logical time
- Property:
  - If  $a \longrightarrow b$ , then  $c(a) < c(b)$
  - If  $a \parallel b$ , then no guarantees

# Lamport clocks: Example



# Concurrency and Lamport clocks

- If  $e1 \longrightarrow e2$ 
  - Then no Lamport clock  $C$  exists with  $C(e1) == C(e2)$

# Concurrency and Lamport clocks

- If  $e1 \longrightarrow e2$ 
  - Then no Lamport clock  $C$  exists with  $C(e1) == C(e2)$
- If  $e1 \parallel e2$ , then there exists a Lamport clock  $C$  such that  $C(e1) == C(e2)$



# The Purpose of Lamport Clocks

- If  $a \longrightarrow b$ , then  $c(a) < c(b)$
- If we order all events by their Lamport clock times
  - We get a partial order, since some events have same time
  - The partial order satisfies “causal relations”

# The purpose of Lamport clocks

- Suppose there are events in different machines
  - Transactions, money in/out, file read, write, copy
- An ordering of events that guarantees preserving causality

# Vector Clocks


- We want a clock such that:
  - If  $a \longrightarrow b$ , then  $c(a) < c(b)$
  - AND
  - If  $c(a) < c(b)$ , then  $a \longrightarrow b$
- Ref: Coulouris et al., V. Garg

# Vector Clocks


- Each process  $i$  maintains a vector  $V_i$
- $V_i$  has  $n$  elements
  - keeps clock  $V_i[j]$  for every other process  $j$
  - On every local event:  $V_i[i] = V_i[i] + 1$
  - On sending a message,  $i$  sends entire  $V_i$
  - On receiving a message at process  $j$ :
    - Takes max element by element
    - $V_j[k] = \max(V_j[k], V_i[k])$ , for  $k = 1, 2, \dots, n$
    - And adds 1 to  $V_j[j]$

# Example

Time



# Another Example



# Comparing Timestamps

- $V = V'$  iff  $V[i] == V'[i]$  for  $i=1,2,\dots,n$
- $V < V'$  iff  $V[i] < V'[i]$  for  $i=1,2,\dots,n$

# Comparing Timestamps

- $V = V'$  iff  $V[i] == V'[i]$  for  $i=1,2,\dots,n$
- $V < V'$  iff  $V[i] < V'[i]$  for  $i=1,2,\dots,n$
- For events  $a, b$  and vector clock  $V$ 
  - $a \longrightarrow b$  iff  $V(a) < V(b)$
- Is this a total order?



# Comparing Timestamps

- $V = V'$  iff  $V[i] == V'[i]$  for  $i=1,2,\dots,n$
- $V \leq V'$  iff  $V[i] \leq V'[i]$  for  $i=1,2,\dots,n$
- For events  $a, b$  and vector clock  $V$ 
  - $a \longrightarrow b$  iff  $V(a) \leq V(b)$
- Two events are concurrent if
  - Neither  $V(a) \leq V(b)$  nor  $V(b) \leq V(a)$

# Vector Clock Examples

- $(1,2,1) \leq (3,2,1)$  but  $(1,2,1) \not\leq (3,1,2)$
- Also  $(3,1,2) \not\leq (1,2,1)$
- No ordering exists

# Vector Clocks

- What are the drawbacks?
- What is the communication complexity?

# Vector Clocks

- What are the drawbacks?
  - Entire vector is sent with message
  - All vector elements ( $n$ ) have to be checked on every message
- What is the communication complexity?
  - $\Omega(n)$  *per message*
  - Increases with time

# Logical Clocks

- There is no way to have perfect knowledge on ordering of events
  - A “true” ordering may not exist..
  - Logical and vector clocks give us a way to have ordering consistent with causality

# Logical time

- Capture just the “happens before” relationship between events
  - Discard the infinitesimal granularity of time
  - Corresponds roughly to causality
- Time at each process is well-defined
  - Definition ( $\rightarrow_i$ ): We say  $e \rightarrow_i e'$  if  $e$  happens before  $e'$  at process  $i$

# Important Points

- Physical Clocks
  - Can keep closely synchronized, but never perfect
- Logical Clocks
  - Encode causality relationship
  - Lamport clocks provide only one-way encoding
  - Vector clocks provide exact causality information